

Overview of outputs of R markdown – Session 15- Assignment

- Predict the no of comments in next H hrs
- Use regression technique
- Report the training accuracy and test accuracy

Feature-Training log regression

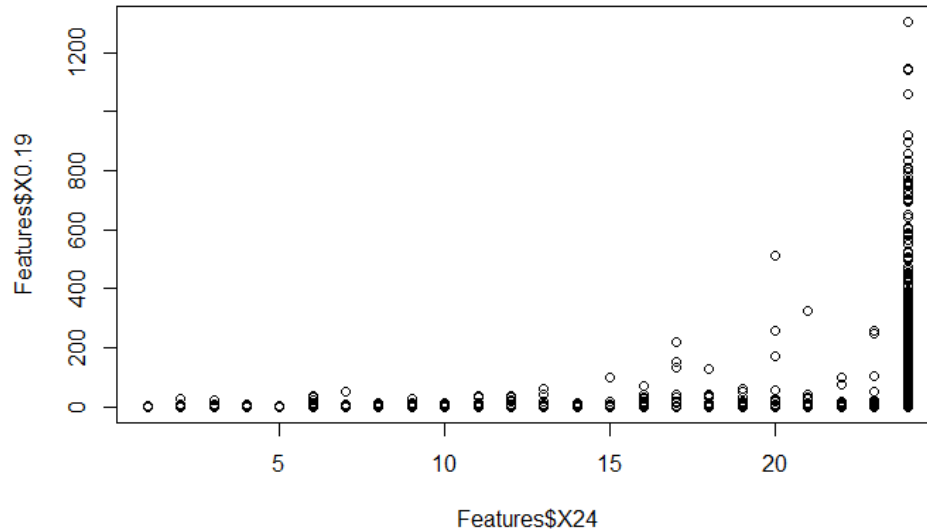
Over view

Attribute Information:

(39 - This describes the H hrs, for which we have the target variable/ comments received.

54 -Target Variable- Decimal Target The no of comments in next H hrs(H is given in Feature no 39).

plot(Features\$X24,Features\$X0.19)



The coding and the main output are given below

```
Features_Train<-Features_norm[1:28003,]  
Features_Test<-Features_norm[28004:40948,]  
library(neuralnet)
```

```
model<-glm(X0.19~X634995+X0+X463+X1+X0.0+X806.0+X11.291044776119403+X1.0+X70.49513846124168+X0.0.1+X806.0.1+X7.574626865671642+X0.0.2+X69.435826365571+X0.0.3+X76.0+X2.6044776119402986+X0.0.4+X8.50550186882253+X0.0.5+X806.0.2+X10.649253731343284+X1.0.1+X70.25478763764251+X.69.0+X806.0.3+X4.970149253731344,data = Features_Train, family = binomial)
```

```
## Warning in eval(family$initialize): non-integer #successes in a binomial
## glm!
```

```
model
```

```
##
```

```
## Call:  glm(formula = X0.19 ~ X634995 + X0 + X463 + X1 + X0.0 + X806.0 +
##      X11.291044776119403 + X1.0 + X70.49513846124168 + X0.0.1 +
##      X806.0.1 + X7.574626865671642 + X0.0.2 + X69.435826365571 +
##      X0.0.3 + X76.0 + X2.6044776119402986 + X0.0.4 + X8.50550186882253 +
##      X0.0.5 + X806.0.2 + X10.649253731343284 + X1.0.1 + X70.25478763764251
+
##      X.69.0 + X806.0.3 + X4.970149253731344, family = binomial,
##      data = Features_Train)
```

```
##
```

```
## Coefficients:
```

(Intercept)	X634995	X0
-8.1890	15.1657	0.6445
X463	X1	X0.0
-23.2171	-0.7102	-26.2338
X806.0	X11.291044776119403	X1.0
2.0551	157.0660	4.4092
X70.49513846124168	X0.0.1	X806.0.1
-8.4151	-10.0358	-6.5078
X7.574626865671642	X0.0.2	X69.435826365571
10.7507	12.4957	3.7177
X0.0.3	X76.0	X2.6044776119402986
-11.5462	4.2625	1.7392
X0.0.4	X8.50550186882253	X0.0.5
-4.4118	-6.5585	18.3120
X806.0.2	X10.649253731343284	X1.0.1
-1.0349	-113.9644	-23.7730
X70.25478763764251	X.69.0	X806.0.3
8.9179	3.4798	5.6877
X4.970149253731344		
NA		

```
##
```

```
## Degrees of Freedom: 28002 Total (i.e. Null); 27976 Residual
```

```
## Null Deviance: 734.5
```

```
## Residual Deviance: 505.6 AIC: 498.2
```

```
summary(model)
```

```
##
```

```
## Call:
```

```

## glm(formula = X0.19 ~ X634995 + X0 + X463 + X1 + X0.0 + X806.0 +
##      X11.291044776119403 + X1.0 + X70.49513846124168 + X0.0.1 +
##      X806.0.1 + X7.574626865671642 + X0.0.2 + X69.435826365571 +
##      X0.0.3 + X76.0 + X2.6044776119402986 + X0.0.4 + X8.50550186882253 +
##      X0.0.5 + X806.0.2 + X10.649253731343284 + X1.0.1 + X70.25478763764251
+
##      X.69.0 + X806.0.3 + X4.970149253731344, family = binomial,
##      data = Features_Train)
##
## Deviance Residuals:
##      Min        1Q      Median        3Q        Max
## -0.59736  -0.07115  -0.05908  -0.02997   2.59350
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -8.1890     2.0025  -4.089 4.33e-05 ***
## X634995         15.1657     7.1199   2.130 0.03317 *
## X0              0.6445     0.7299   0.883 0.37721
## X463           -23.2171     8.4336  -2.753 0.00591 **
## X1             -0.7102     0.6881  -1.032 0.30200
## X0.0           -26.2338    70.6481  -0.371 0.71039
## X806.0          2.0551     4.4873   0.458 0.64696
## X11.291044776119403 157.0660   110.2234   1.425 0.15416
## X1.0            4.4092     47.7658   0.092 0.92645
## X70.49513846124168  -8.4151    20.2766  -0.415 0.67813
## X0.0.1         -10.0358    23.2964  -0.431 0.66662
## X806.0.1        -6.5078     4.5938  -1.417 0.15658
## X7.574626865671642  10.7507    35.1815   0.306 0.75993
## X0.0.2          12.4957    18.7866   0.665 0.50596
## X69.435826365571    3.7177     4.9712   0.748 0.45455
## X0.0.3         -11.5462    19.7890  -0.583 0.55958
## X76.0            4.2625     2.3504   1.813 0.06976 .
## X2.6044776119402986  1.7392     6.1798   0.281 0.77838
## X0.0.4          -4.4118     5.1254  -0.861 0.38937
## X8.50550186882253  -6.5585     5.5094  -1.190 0.23388
## X0.0.5          18.3120    64.8921   0.282 0.77780
## X806.0.2        -1.0349     4.4386  -0.233 0.81564
## X10.649253731343284 -113.9644    86.4334  -1.319 0.18733
## X1.0.1          -23.7730    34.6937  -0.685 0.49320
## X70.25478763764251   8.9179    19.8051   0.450 0.65250
## X.69.0           3.4798     4.0256   0.864 0.38736
## X806.0.3         5.6877     4.8682   1.168 0.24267
## X4.970149253731344      NA         NA      NA      NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 734.54  on 28002  degrees of freedom
## Residual deviance: 505.64  on 27976  degrees of freedom

```

```
## AIC: 498.16
##
## Number of Fisher Scoring iterations: 9

predict<-predict(model)
head(predict)

##          1          2          3          4          5          6
## -5.584117 -5.584117 -5.584117 -5.584117 -5.584117 -5.584117
```

```
fit <- glm(X0.19~X634995+X0+X463+X1+X0.0+X806.0+X11.291044776119403+X1.0+X70.49513846124168+X0.0.1+X806.0.1+X7.574626865671642+X0.0.2+X69.435826365571+X0.0.3+X76.0+X2.6044776119402986+X0.0.4+X8.50550186882253+X0.0.5+X806.0.2+X10.649253731343284+X1.0.1+X70.25478763764251+X.69.0+X806.0.3+X4.970149253731344,data = Features_Train, family = binomial(link='logit'))
```

```
## Warning in eval(family$initialize): non-integer #successes in a binomial
## glm!
```

```
summary(fit)
```

```
##
## Call:
## glm(formula = X0.19 ~ X634995 + X0 + X463 + X1 + X0.0 + X806.0 +
##      X11.291044776119403 + X1.0 + X70.49513846124168 + X0.0.1 +
##      X806.0.1 + X7.574626865671642 + X0.0.2 + X69.435826365571 +
##      X0.0.3 + X76.0 + X2.6044776119402986 + X0.0.4 + X8.50550186882253 +
##      X0.0.5 + X806.0.2 + X10.649253731343284 + X1.0.1 + X70.25478763764251
##      +
##      X.69.0 + X806.0.3 + X4.970149253731344, family = binomial(link = "logi
##      t"),
##      data = Features_Train)
##
```

```
## Deviance Residuals:
```

```
##      Min      1Q   Median      3Q      Max
## -0.59736 -0.07115 -0.05908 -0.02997  2.59350
##
```

```
## Coefficients: (1 not defined because of singularities)
```

```
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -8.1890     2.0025  -4.089 4.33e-05 ***
## X634995         15.1657     7.1199   2.130  0.03317 *
## X0              0.6445     0.7299   0.883  0.37721
## X463           -23.2171     8.4336  -2.753  0.00591 **
## X1             -0.7102     0.6881  -1.032  0.30200
## X0.0           -26.2338    70.6481  -0.371  0.71039
```

```

## X806.0      2.0551      4.4873      0.458      0.64696
## X11.291044776119403 157.0660 110.2234      1.425      0.15416
## X1.0         4.4092      47.7658      0.092      0.92645
## X70.49513846124168 -8.4151      20.2766     -0.415      0.67813
## X0.0.1       -10.0358     23.2964     -0.431      0.66662
## X806.0.1      -6.5078      4.5938     -1.417      0.15658
## X7.574626865671642 10.7507      35.1815      0.306      0.75993
## X0.0.2        12.4957     18.7866      0.665      0.50596
## X69.435826365571    3.7177      4.9712      0.748      0.45455
## X0.0.3       -11.5462     19.7890     -0.583      0.55958
## X76.0         4.2625      2.3504      1.813      0.06976 .
## X2.6044776119402986  1.7392      6.1798      0.281      0.77838
## X0.0.4       -4.4118      5.1254     -0.861      0.38937
## X8.50550186882253   -6.5585      5.5094     -1.190      0.23388
## X0.0.5        18.3120     64.8921      0.282      0.77780
## X806.0.2      -1.0349      4.4386     -0.233      0.81564
## X10.649253731343284 -113.9644     86.4334     -1.319      0.18733
## X1.0.1       -23.7730     34.6937     -0.685      0.49320
## X70.25478763764251   8.9179     19.8051      0.450      0.65250
## X.69.0         3.4798      4.0256      0.864      0.38736
## X806.0.3       5.6877      4.8682      1.168      0.24267
## X4.970149253731344      NA          NA          NA          NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 734.54  on 28002  degrees of freedom
## Residual deviance: 505.64  on 27976  degrees of freedom
## AIC: 498.16
##
## Number of Fisher Scoring iterations: 9

library(ResourceSelection)


## ResourceSelection 0.3-2    2017-02-28

hoslem.test(Features_Train$X0.19, fitted(fit))


##
## Hosmer and Lemeshow goodness of fit (GOF) test
##
## data:  Features_Train$X0.19, fitted(fit)
## X-squared = 14.466, df = 8, p-value = 0.07041

#plot the fitted model
plot(fit$fitted.values)


```



```
x <- 1:20  
y <- 21:40  
library(MASS)  
boxcox(y~x)
```




```
plot(1/y^2~x)
```



```
g <- roc(X0.19~ pred, data = Features_Test)
```



```
# summarize results
#confusionMatrix<- confusionMatrix(pred$predictions,pred$X0.19)
#confusionMatrix
barchart(Features_Test$pred_X0.19)
```



Neural Network model prediction

```
library(readr)

dataset <- read_delim("Dataset.zip", ";", escape_double = FALSE, trim_ws = TRUE)

View(Features_TestSet)

dim(Features_TestSet)

head(Features_TestSet)

str(Features_TestSet)

names(Features_TestSet)
```

```

dim(Features_Variant_1)
Features<-Features_Variant_1
names(Features)
head(Features)
head(Test_Case_1)
normalize<-function(x) {return((x-min(x))/(max(x)-min(x)))}
Features_norm<-as.data.frame(lapply(Features,normalize))
Features_norm
names(Features_norm)
Features_Train<-Features_norm[1:28003,]
Features_Test<-Features_norm[28004:40948,]
library(neuralnet)
set.seed(123)
Features_model<-neuralnet(X0_19~
X634995+X0+X463+X1+X0.0+X806.0+X11.291044776119403+X1.0+X70.49513846124168+X0.0_1+X806.
0_1+X7.574626865671642+X0.0_2+X69.435826365571+X0.0_3+X76.0+X2.6044776119402986+X0.0_4+
X8.50550186882253+X0.0_5+X806.0_2+X10.649253731343284+X1.0_1+X70.25478763764251+X.69.0+
X806.0_3+X4.970149253731344+X0.0_6+X69.85058043098057+X0_1+X0_2+X0_3+X0_4+X0_5+X65+X1
66+X2,data = Features_Train)

plot(Features_model)

model_results<-compute(Features_model,Features_Test[1:37])

predicted_X0_19<-model_results$net.result

nrow(predicted_X0_19)

head(predicted_X0_19)

round(cor( predicted_X0_19,Features_Test$X0_19),2)


```

```

[1] 12945
      [,1]
28004 0.001069433699
28005 0.001070214932
28006 0.001087235290
28007 0.001069451002
28008 0.001589321488
28009 0.001273782979
      [,1]
[1,] 0.65

```

The correlation (0.65) is the highest than the others given below. Higher is better.




```
set.seed(123)
```

```
Features_model<-
```


```
neuralnet(X0_19~X634995+X0+X463+X1+X0.0+X806.0+X11.291044776119403+X1.0+X70.49513846124
168+X0.0_1+X806.0_1+X7.574626865671642+X0.0_2+X69.435826365571+X0.0_3+X76.0+X2.60447761
19402986+X0.0_4+X8.50550186882253+X0.0_5+X806.0_2+X10.649253731343284+X1.0_1+X70.254787
63764251+X.69.0+X806.0_3+X4.970149253731344,data = Features_Train)
```

```
plot(Features_model)
```

Model1



Model2



```
model_results<-compute(Features_model,Features_Test[1:27])
```

```
predicted_X0_19<-model_results$net.result
```

```
nrow(predicted_X0_19)
```

```
head(predicted_X0_19)
```

```
round(cor( predicted_X0_19,Features_Test$X0_19),2)
```

```
[1] 12945
      [,1]
28004 0.002802219847
28005 0.002802219847
28006 0.002802219847
28007 0.002802219847
28008 0.002802219847
28009 0.002802219847
      [,1]
[1,] 0.34
```


```
set.seed(123)
```

```
Features_model<-
```


```
neuralnet(X0.19~X634995+X0+X463+X1+X0.0+X806.0+X11.291044776119403+X1.0+X70.495138461241  
68+X0.0.1+X806.0.1+X7.574626865671642,data = Features_Train)
```

```
plot(Features_model)
```

Model1



Model2



Model1

```
[1] 12945
      [,1]
28004 0.002660508374
28005 0.002660508374
28006 0.002660508374
28007 0.002660508374
28008 0.002660508374
28009 0.002660508374
      [,1]
[1,] 0.34
```

Model2

```
[1] 12945
      [,1]
28004 0.001997766546
28005 0.001997766546
28006 0.001997766546
28007 0.001997766546
28008 0.001997766546
28009 0.001997766546
      [,1]
[1,] 0.32
```

non-integer #successes in a binomial glm!

```
Call: glm(formula = x0.19 ~ x634995 + x0 + x463 + x1 + x0.0 + x806.0 +  
  x11.291044776119403 + x1.0 + x70.49513846124168 + x0.0.1 +  
  x806.0.1 + x7.574626865671642 + x0.0.2 + x69.435826365571 +  
  x0.0.3 + x76.0 + x2.6044776119402986 + x0.0.4 + x8.50550186882253 +  
  x0.0.5 + x806.0.2 + x10.649253731343284 + x1.0.1 + x70.25478763764251 +  
  x.69.0 + x806.0.3 + x4.970149253731344, family = binomial,  
  data = Features_Train)
```

Coefficients:

(Intercept)	x634995	x0	
x463			
-8.1889873	15.1656626	0.6445092	-
23.2171336			
x1	x0.0	x806.0	
x11.291044776119403			
-0.7102152	-26.2338417	2.0551480	
157.0660451			
x1.0	x70.49513846124168	x0.0.1	
x806.0.1			
4.4092485	-8.4150733	-10.0357869	-
6.5077822			
x7.574626865671642	x0.0.2	x69.435826365571	
x0.0.3			
10.7507051	12.4957414	3.7176835	-
11.5461700			
x76.0	x2.6044776119402986	x0.0.4	
x8.50550186882253			
4.2624506	1.7391687	-4.4118131	-
6.5584531			
x0.0.5	x806.0.2	x10.649253731343284	
x1.0.1			
18.3120490	-1.0348589	-113.9643962	-
23.7729803			
x70.25478763764251	x.69.0	x806.0.3	
x4.970149253731344			
8.9179436	3.4798448	5.6876570	
NA			

Degrees of Freedom: 28002 Total (i.e. Null); 27976 Residual

Null Deviance: 734.5361

Residual Deviance: 505.6398 AIC: 498.1562